

# ECON 5140: Applied Econometrics

## Homework 1: GLMs & Time Series Foundations

---

### PART 1: ANALYSIS PROBLEMS

#### Problem 1: Logistic Regression - Loan Default Prediction

A bank wants to predict loan defaults based on credit score. They have the following data from a logistic regression:

Model:  $\log(p/(1-p)) = \beta_0 + \beta_1(\text{CreditScore})$

Estimated coefficients:

- $\beta_0 = 5.2$
- $\beta_1 = -0.008$

where  $p = P(\text{Default} = 1 | \text{CreditScore})$

Questions:

- Calculate the predicted probability of default for someone with a credit score of 650.
  - Calculate the predicted probability of default for someone with a credit score of 750.
  - At what credit score is the predicted probability of default equal to 0.5? (This is the decision boundary)
  - Interpret the coefficient  $\beta_1 = -0.008$  in terms of odds ratios. What happens to the odds of default when credit score increases by 100 points?
-

## Problem 2: Poisson Regression - Website Visits

An e-commerce company models daily website visits using Poisson regression.

Model:  $\log(E[\text{Visits}|\mathbf{X}]) = \beta_0 + \beta_1(\text{Ad\_Spend}) + \beta_2(\text{Weekend})$

where:

- Ad\_Spend is in thousands of dollars
- Weekend = 1 if weekend, 0 otherwise

Estimated coefficients:

- $\beta_0 = 5.5$
- $\beta_1 = 0.12$
- $\beta_2 = 0.30$

Questions:

- Predict the expected number of visits on a weekday with \$5,000 ad spend.
- Interpret  $\beta_1 = 0.12$ . What is the percentage change in expected visits for each additional \$1,000 in ad spending?
- Interpret  $\beta_2 = 0.30$ . What is the percentage increase in visits on weekends compared to weekdays?
- If the company wants to achieve 500 visits on a weekday, how much should they spend on ads?

---

## Problem 3: Autocorrelation Calculation

A tech company tracks daily active users (in thousands). Here are 10 consecutive days of data:

Day	1	2	3	4	5	6	7	8	9	10
Users ( $Y_t$ )	45	48	50	49	52	54	53	56	58	57

Questions:

- a) Calculate the mean ( $\mu$ ) and standard deviation ( $\sigma$ ) of the series.
- b) Calculate the lag-1 autocorrelation  $\rho(1) = \text{Corr}(Y_t, Y_{t-1})$ :
- Hint: Create two series:  $Y_1$  to  $Y_9$  and  $Y_2$  to  $Y_{10}$ , then calculate correlation
- c) What does this autocorrelation value tell you about the persistence in daily active users?
- d) Based on this autocorrelation, would knowing Day 10's value help forecast Day 11? Why?
- 

## Problem 4: Moving Averages

Daily website traffic (in thousands of visits):

<b>Day</b>	<b>1</b>	<b>2</b>	<b>3</b>	<b>4</b>	<b>5</b>	<b>6</b>	<b>7</b>
Visit s	10	12	11	15	13	16	14

Questions:

- a) Calculate a 3-day centered moving average for Day 4:
- $\text{MA}_3(4) = (Y_3 + Y_4 + Y_5) / 3$
- b) Calculate a 3-day trailing moving average for Day 5:
- $\text{MA}_3(5) = (Y_3 + Y_4 + Y_5) / 3$
- c) For Day 4, calculate a weighted moving average with weights (0.25, 0.50, 0.25):
- $\text{WMA}(4) = 0.25 \times Y_3 + 0.50 \times Y_4 + 0.25 \times Y_5$
- d) Why might the weighted average produce a smoother trend estimate?

# PART 2: CODING PROBLEMS

## Problem 5: Customer Purchase Prediction & Time Series Analysis

You will work with two datasets: customer purchase behavior (GLM) and e-commerce sales over time (Time Series).

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy import stats
import statsmodels.api as sm
from statsmodels.discrete.discrete_model import Logit, Probit, Poisson
from statsmodels.tsa.seasonal import STL, seasonal_decompose
from statsmodels.graphics.tsaplots import plot_acf
import warnings
warnings.filterwarnings('ignore')

# Set random seed for reproducibility
np.random.seed(42)

print("=" * 70)
print("ECON 5140 - HOMEWORK 1")
print("Part A: Generalized Linear Models")
print("Part B: Time Series Decomposition")
print("=" * 70)

# =====
# DATASET 1: CUSTOMER PURCHASE DATA (for GLM analysis)
# =====
print("\n" + "=" * 70)
print("DATASET 1: Customer Purchase Behavior")
print("=" * 70)

n_customers = 1000

# Generate customer features
age = np.random.normal(35, 10, n_customers)
income = np.random.normal(50, 15, n_customers) # in thousands
time_on_site = np.random.gamma(2, 3, n_customers) # in minutes

# True relationship (latent variable)
z = -3 + 0.05*age + 0.04*income + 0.15*time_on_site + np.random.normal(0, 1, n_customers)
```

```

# Generate binary outcome (Purchase: 1=Yes, 0=No)
purchase = (z > o).astype(int)

# Create DataFrame
df_customers = pd.DataFrame({
    'Age': age,
    'Income': income,
    'TimeOnSite': time_on_site,
    'Purchase': purchase
})

print(f"Number of customers: {len(df_customers)}")
print(f"Purchase rate: {df_customers['Purchase'].mean():.2%}")
print(f"\nFirst 5 rows:")
print(df_customers.head())

# =====
# DATASET 2: E-COMMERCE SALES TIME SERIES
# =====
print("\n" + "=" * 70)
print("DATASET 2: E-commerce Daily Sales")
print("=" * 70)

# Create 2 years of daily data
dates = pd.date_range('2024-01-01', '2025-12-31', freq='D')
n_days = len(dates)
t = np.arange(n_days)

# Components
trend = 1000 + 2*t + 0.01*t**2
yearly_seasonal = 200 * np.sin(2*np.pi*t/365) + 150 * np.cos(2*np.pi*t/365)
weekly_seasonal = 100 * np.sin(2*np.pi*t/7)

# Special events
special_events = np.zeros(n_days)
for year in [2024, 2025]:
    # Black Friday
    bf_date = pd.Timestamp(f'{year}-11-24')
    bf_idx = (dates == bf_date)
    special_events[bf_idx] = 800

    # Christmas
    xmas_idx = (dates >= f'{year}-12-20') & (dates <= f'{year}-12-25')
    special_events[xmas_idx] = 400

# Random noise
noise = np.random.normal(0, 50, n_days)

# Combine components

```

```

sales = trend + yearly_seasonal + weekly_seasonal + special_events + noise
sales = np.maximum(sales, 0)

# Create DataFrame
df_sales = pd.DataFrame({
    'Date': dates,
    'Sales': sales,
    'DayOfWeek': dates.dayofweek,
    'Month': dates.month,
    'IsWeekend': dates.dayofweek >= 5
})
df_sales.set_index('Date', inplace=True)

print(f'Date range: {df_sales.index[0].date()} to {df_sales.index[-1].date()}')
print(f'Number of days: {len(df_sales)}')
print(f'\nSales Statistics:')
print(df_sales['Sales'].describe())

# =====
# PART A: GENERALIZED LINEAR MODELS
# =====
print("\n" + "=" * 70)
print("PART A: GENERALIZED LINEAR MODELS")
print("=" * 70)

# -----
# A1: Exploratory Data Analysis (GLM)
# -----
print("\n" + "-" * 70)
print("A1: Exploratory Data Analysis")
print("-" * 70)

# YOUR CODE:
# 1. Create box plots comparing Age, Income, and TimeOnSite
#   between purchasers and non-purchasers
#   - Use 3 subplots (1 row, 3 columns)
#
# 2. Calculate and print mean values for each group:
#   - Mean Age: Purchasers vs Non-purchasers
#   - Mean Income: Purchasers vs Non-purchasers
#   - Mean TimeOnSite: Purchasers vs Non-purchasers
#
# 3. Create a correlation matrix heatmap for the features

# -----
# A2: Linear Probability Model (LPM)
# -----
print("\n" + "-" * 70)
print("A2: Linear Probability Model")

```

```
print("-" * 70)

# YOUR CODE:
# 1. Fit OLS model: Purchase ~ Age + Income + TimeOnSite
#   - Add constant using sm.add_constant()
#   - Use sm.OLS()
#
# 2. Print regression summary
#
# 3. Calculate predicted probabilities
#   - Count how many predictions are outside [0, 1]
#   - Print the percentage of invalid predictions
#
# 4. Create histogram of predicted probabilities
#   - Mark the [0, 1] boundaries with vertical lines
```

```
# -----
# A3: Logistic Regression
# -----
print("\n" + "-" * 70)
print("A3: Logistic Regression")
print("-" * 70)
```

```
# YOUR CODE:
# 1. Fit logistic regression: sm.Logit()
#
# 2. Print summary and extract:
#   - Coefficients
#   - Odds ratios: np.exp(coefficients)
#   - p-values
#
# 3. Interpret each coefficient:
#   - Age: Effect on log-odds
#   - Income: Effect on log-odds
#   - TimeOnSite: Effect on log-odds
#
# 4. Calculate predicted probabilities
#   - Verify all are in [0, 1]
#   - Create histogram
```

```
# -----
# A4: Prediction for New Customers
# -----
print("\n" + "-" * 70)
print("A4: Predictions for New Customers")
print("-" * 70)
```

```

# New customers
new_customers = pd.DataFrame({
    'Age': [25, 35, 45, 55],
    'Income': [30, 50, 70, 90],
    'TimeOnSite': [2, 5, 8, 10]
})

# YOUR CODE:
# 1. Use your logistic model to predict purchase probability
#   for each new customer
#
# 2. Create a nice formatted table showing:
#   - Customer features
#   - Predicted probability
#   - Classification (Purchase = 1 if p > 0.5)
#
# 3. Which customer is most likely to purchase? Why?

# =====
# PART B: TIME SERIES ANALYSIS
# =====
print("\n" + "=" * 70)
print("PART B: TIME SERIES ANALYSIS")
print("=" * 70)

# -----
# B1: Time Series Visualization
# -----
print("\n" + "-" * 70)
print("B1: Time Series Visualization")
print("-" * 70)

# YOUR CODE:
# 1. Create time series plot of daily sales
#   - Full 2-year series
#   - Label axes appropriately
#
# 2. Create seasonal subseries plots:
#   - Box plot: Sales by day of week
#   - Box plot: Sales by month
#
# 3. Calculate and print:
#   - Mean sales by day of week
#   - Mean sales by month
#
# 4. What patterns do you observe?

# -----
# B2: Stationarity Assessment

```

```

# -----
print("\n" + "-" * 70)
print("B2: Stationarity Check")
print("-" * 70)

# YOUR CODE:
# 1. Create a plot with 3 subplots:
#   - Original series
#   - Rolling mean (30-day window)
#   - Rolling std (30-day window)
#
# 2. Is the series stationary? Justify based on:
#   - Does mean change over time?
#   - Does variance change over time?
#
# 3. Compare first 6 months vs last 6 months:
#   - Calculate mean and std for each period
#   - Print comparison

# -----
# B3: Autocorrelation Analysis
# -----
print("\n" + "-" * 70)
print("B3: Autocorrelation Function")
print("-" * 70)

# YOUR CODE:
# 1. Plot ACF for up to 60 lags
#   - Use plot_acf from statsmodels
#
# 2. Calculate specific autocorrelations manually:
#   - Lag 1 (yesterday)
#   - Lag 7 (last week)
#   - Lag 30 (last month)
#   Use: np.corrcoef(sales[:-lag], sales[lag:])[0,1]
#
# 3. Interpret:
#   - Do you see weekly patterns?
#   - How persistent is the autocorrelation?

# -----
# B4: STL Decomposition
# -----
print("\n" + "-" * 70)
print("B4: STL Decomposition")
print("-" * 70)

# YOUR CODE:
# 1. Apply STL decomposition with weekly seasonality:

```

```

#   stl = STL(df_sales['Sales'], seasonal=7, robust=True)
#   result = stl.fit()
#
# 2. Plot all four components:
#   - Observed
#   - Trend
#   - Seasonal (weekly pattern)
#   - Remainder
#
# 3. Analyze each component:
#   - What is the trend pattern?
#   - What is the weekly seasonal pattern?
#   - Are special events visible in remainder?
#
# -----
# B5: Remainder Diagnostics
# -----
print("\n" + "-" * 70)
print("B5: Remainder Analysis")
print("-" * 70)

# YOUR CODE:
# 1. Extract remainder from STL
#
# 2. Create diagnostic plots:
#   - Time series plot of remainder
#   - Histogram of remainder
#   - ACF of remainder
#
# 3. Statistical tests:
#   - Mean (should be ≈ 0)
#   - Standard deviation
#   - Check normality
#
# 4. Identify outliers:
#   - Find days where |remainder| > 3×std
#   - Print dates and investigate
#   - Are Black Friday and Christmas visible?

```